

6.034: Lecture 5

Adversarial search & games

Can computers play chess?

Professor Robert C. Berwick

Chess (Yes, Chess) Is Now a Streaming Obsession

Hikaru Nakamura

Viewers are flocking to games during the pandemic, entranced by a charismatic grandmaster and his lightning-fast play.





AI Methods

■ **Problem solving**

- *G+T, search, optimal search, games, constraint satisfaction*

□ **Inference**

- *rule-based systems, Bayesian inference*

□ **Machine learning**

- *k-nearest neighbors, id trees, neural nets, deep neural nets, support vector machines, genetic algorithms, near miss/one-shot*

□ **Communication, perception, action**

- *natural language processing, vision, robotics*

Menu for today

Games: Adversarial search

- ☐ Intelligence & Chess
- ☐ Ways to Play Chess
- ☐ How to search: Minimax algorithm
- ☐ Improving search: Alpha-beta pruning
- ☐ Progressive Deepening
- ☐ Reflections on contemporary chess computers & Gold Star ideas

A very brief history


Philosophical Magazine, Ser.7, Vol. 41, No. 314 - March 1950.

XXII. Programming a Computer for Playing Chess¹

By CLAUDE E. SHANNON

Bell Telephone Laboratories, Inc., Murray Hill, N.J.²

[Received November 8, 1949]



Claude Shannon
(1916-2001)

1. INTRODUCTION

This paper is concerned with the problem of constructing a computing routine or "program" for a modern general purpose computer which will enable it to play chess. Although perhaps of no practical importance, the question is of theoretical interest, and it is hoped that a satisfactory solution of this problem will act as a wedge in attacking other problems of a similar nature and of greater significance. Some possibilities in this direction are: -

The chess machine is an ideal one to start with, since: (1) the problem is sharply defined both in allowed operations (the moves) and in the ultimate goal (checkmate); (2) it is neither so simple as to be trivial nor too difficult for satisfactory solution; (3) chess is generally considered to require "thinking" for skilful play; a solution of this problem will force us either to admit the possibility of a mechanized thinking or to further restrict our concept of "thinking"; (4) the discrete structure of chess fits well into the digital nature of modern computers.

ALCHEMY AND ARTIFICIAL INTELLIGENCE

Hubert L. Dreyfus

December 1965

remains unimproved. Burton Bloom at M.I.T. has made the latest attempt to write a chess program; like all the others, it plays a stupid game. In fact, in the nine years since the Los Alamos program beat a weak player, in spite of a great investment of time, energy, and ink, the only improvement seems to be that a machine now plays poorly on an eight-by-eight rather than a six-by-six board. According to Newell, Shaw, and Simon themselves, evaluating the Los Alamos, the IBM, and the NSS programs: "All three programs play roughly the same quality of chess (mediocre) with roughly the same amount of computing time" [20:14]. Still no chess program can play even amateur chess, and the world championship tournament is only two years away.



The Greenblatt chess program

by *RICHARD D. GREENBLATT,*

DONALD E. EASTLAKE, III,

and

STEPHEN D. CROCKER

*Massachusetts Institute of Technology
Cambridge, Massachusetts*

INTRODUCTION

Since mid-November 1966 a chess program has been under development at the Artificial Intelligence Laboratory of Project MAC at M.I.T. This paper describes the state of the program as of August 1967 and gives some of the details of the heuristics and algorithms employed.

Fall Joint Computer Conference, 1967




Richard Greenblatt

MIT AI Lab MacHack IV

H. Dreyfus vs. MacHack IV

Dreyfus vs. "Robert Q"
1967

Prof. Dreyfus
loses after 38
moves



1.5 Computers Can't Play Chess.

1.5.1 Nor Can Dreyfus.

THE ARTIFICIAL INTELLIGENCE OF HUBERT L. DREYFUS

A Budget of Fallacies

by


Seymour Papert

January 1968.


I hope readers will pardon a digression to get a debating point off my chest. They will surely understand why I find it irresistible.

In Alchemy and Phenomenology Dreyfus discusses the weakness of chess-playing programs. He plainly gives the impression that they can typically be defeated by human novices. His twice recounted story of how a ten-year-old child defeated a program constructed by Newell, Shaw and Simon was gleefully quoted by the New Yorker and other popular magazines as demonstrating the futility of Artificial Intelligence. While Phenomenology was in press I had the pleasure of arranging for Dreyfus to play against Richard Greenblatt's chess program at M.I.T. and seeing him very roundly trounced. The newsletter SIGART reprinted the game with no comment beyond one phrase from Alchemy:

"... no chess program can play even amateur chess." (p. 10)



John McCarthy



Donald Michie


Versus.



Chess Grandmaster David Levy
Bet in 1968, 1250 pounds

How could a computer play chess?

1. Human like : Describe board via pawn structure, King safety, good time to castle....use tactics, strategy then make a move
2. If-then rules – make most plausible move
3. Look ahead and evaluate



$$S = g(f_1, f_2, \dots, f_n)$$


$$= c_1 f_1 + c_2 f_2 + \dots + c_n f_n$$

Linear
Scoring
Polynomial

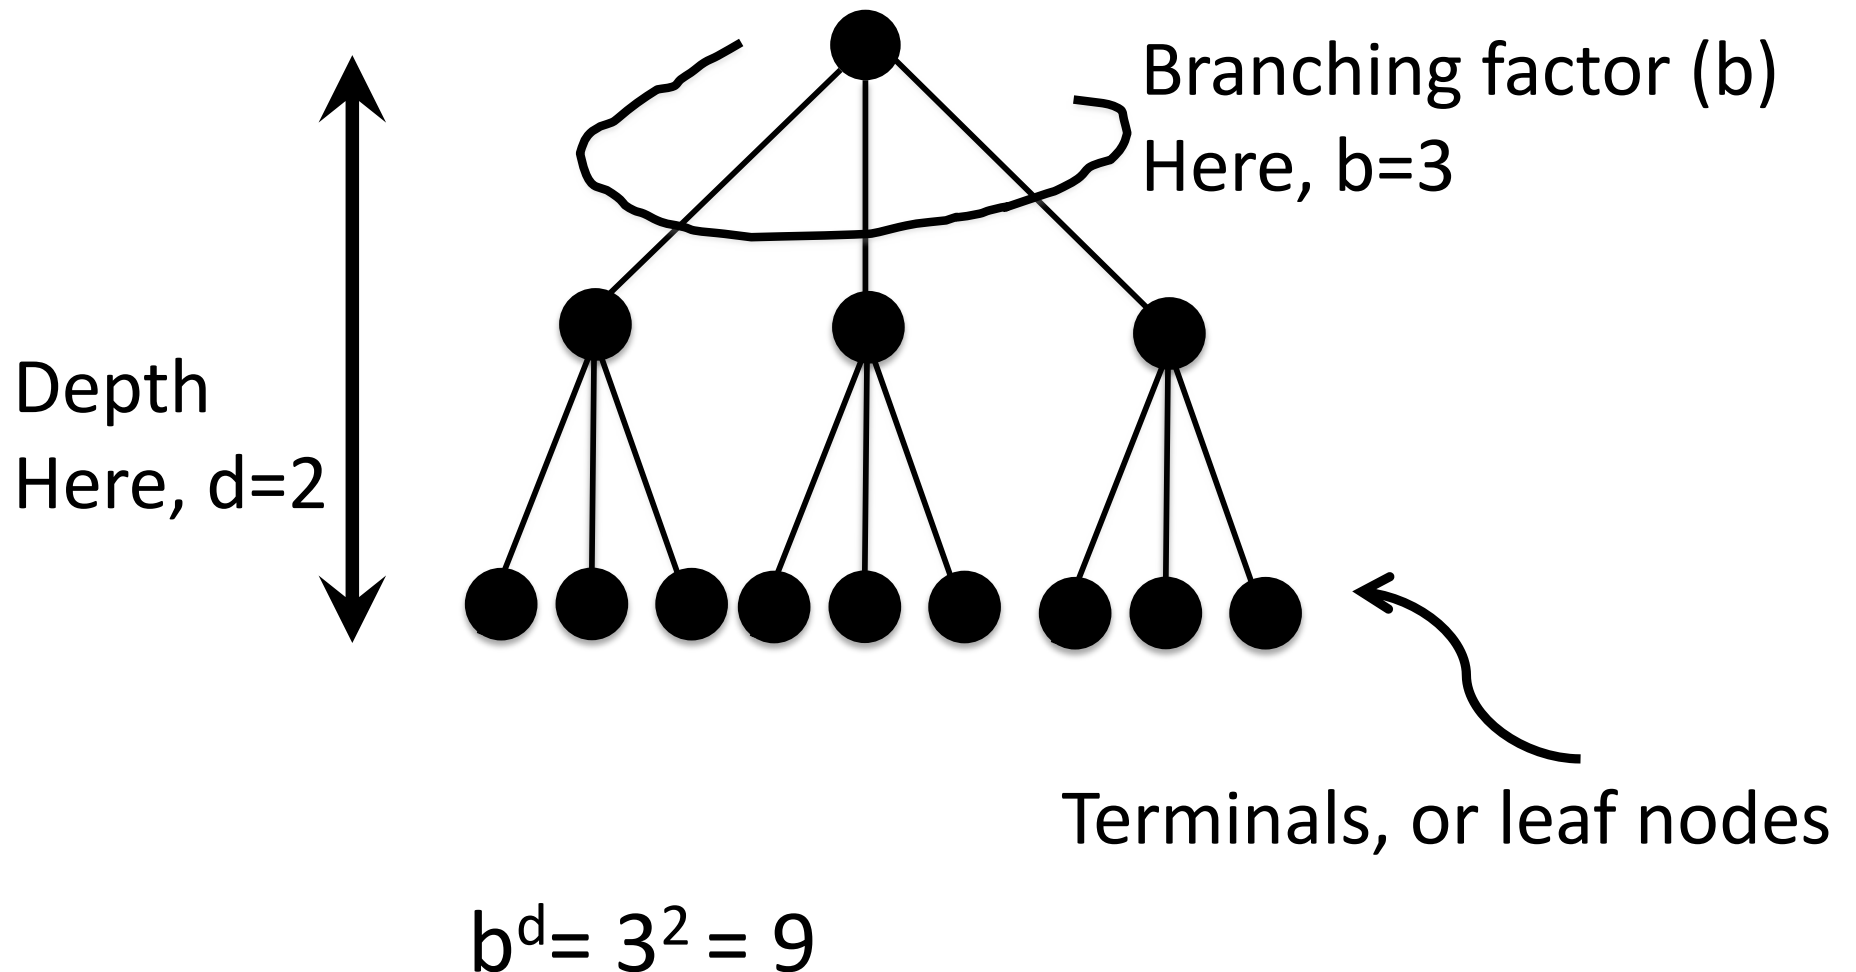
How could a computer play chess?

1. Human like
2. If-then rules: look & make most plausible move
3. Look ahead and evaluate
4. British Museum algorithm (exhaustive search)?

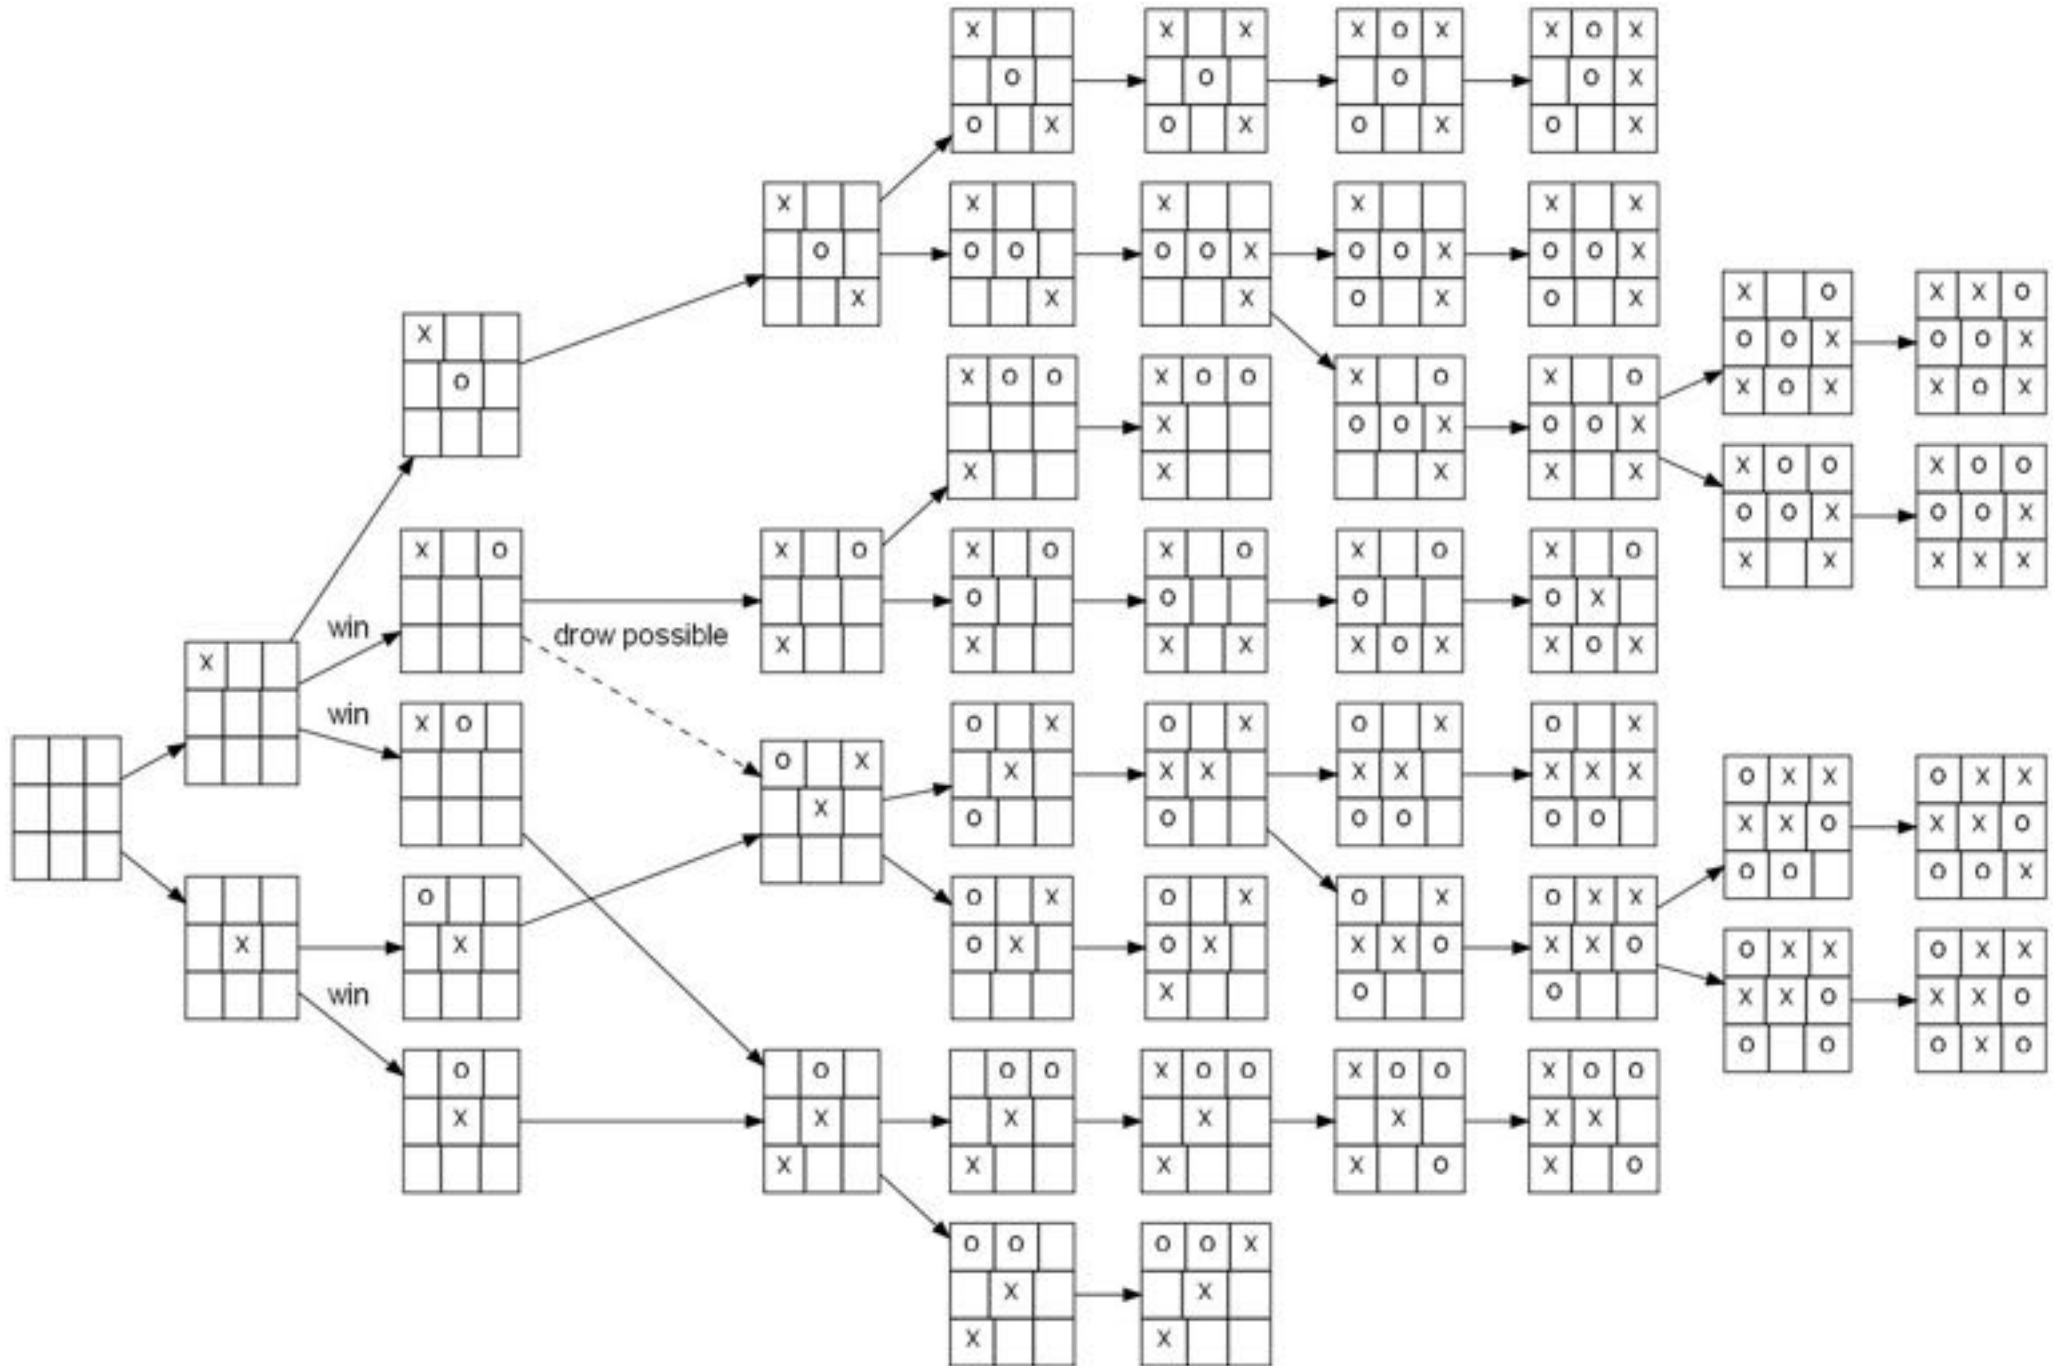
Vocabulary for game trees



Vocabulary for game trees



Tic-tac-toe: complete game tree space



Could we use cloud computing to
evaluate all possible chess moves in the
same way?

10^{120} moves

Could we use cloud computing to evaluate all possible chess moves in the same way?

10^{120} moves

$$\frac{\frac{10^{80} \text{ atoms in universe}}{\pi \times 10^7 \text{ seconds/year}}}{\frac{10^9 \text{ nanoseconds/second}}{10^{10} \text{ years since Big Bang}}}$$


$\approx 10^{26}$ nanoseconds x 10^{80} atoms
 $\approx 10^{106}$ ops, each atom operating
@ nanosecond speeds


Not enough compute power for
 10^{120} moves!

How could a computer play chess?

1. Human like
2. If-then— make most plausible move
3. Look ahead and evaluate
4. British Museum algorithm (exhaustive search)
5. Look ahead as far as possible

Minimax algorithm for searching game tree to find optimal move sequence





Note that the best minimax path is preserved under any transform of the leaf scores that preserves their rank ordering



Games

17:13:35 EDT 02-Aug-2020

- Goal trees
- Search
 - Basic search
 - Optimal search
 - Games**
 - Monte Carlo
- Constraint satisfaction
- Biological mimetics
- Learning
- Neural nets
- Bayes nets
- Miscellaneous

Start New tree Hide choices

Theoretical minimum / actual / maximum: 7 / 16 / 16
Minimax only


MAX

MIN

MAX

MIN

Static
Values



Example

- ☒ Classroom example
- ☐ Quiz example
- ☐ Good news example

Arrangement

- ☐ Random numbers
- ☐ Increasing left to right
- ☐ Decreasing left to right

Search type

- ☒ Mini-max only
- ☐ Mini-max with alpha-beta


Shape

- ☒ 2 x 4
- ☐ 4 x 2
- ☐ 2 x 6
- ☐ 6 x 2

Delay



Minimax with alpha-beta (α - β) pruning



α value = MAX is guaranteed to gain at least this much, or $\geq \alpha$; a floor on MAX's gain.


β value = MIN is guaranteed to lose at most this much, or $\leq \beta$; a ceiling on MIN's loss.

Note: $[\alpha < \beta]$, or floor < ceiling. Why?

Q: Alpha-Beta pruning value = Minimax value w/o pruning?


Ans: 2

A deeper game tree illustrates how much pruning alpha-beta can do ($b = 2$, $d = 3$)




A deeper game tree illustrates how much pruning alpha-beta can do


First half of tree (max- α values; min- β values)



Second half of tree: cutoff whenever $\alpha \geq \beta$



A deeper game tree



6 out of 16 static evaluations are not made



Dead horse principle (principle AWP)



Games

13:22:46 EDT 02-Aug-2020

- Goal trees
- Search
 - Basic search
 - Optimal search
 - Games**
 - Monte Carlo
- Constraint satisfaction
- Biological mimetics
- Learning
- Neural nets
- Bayes nets
- Miscellaneous

Start New tree Hide choices

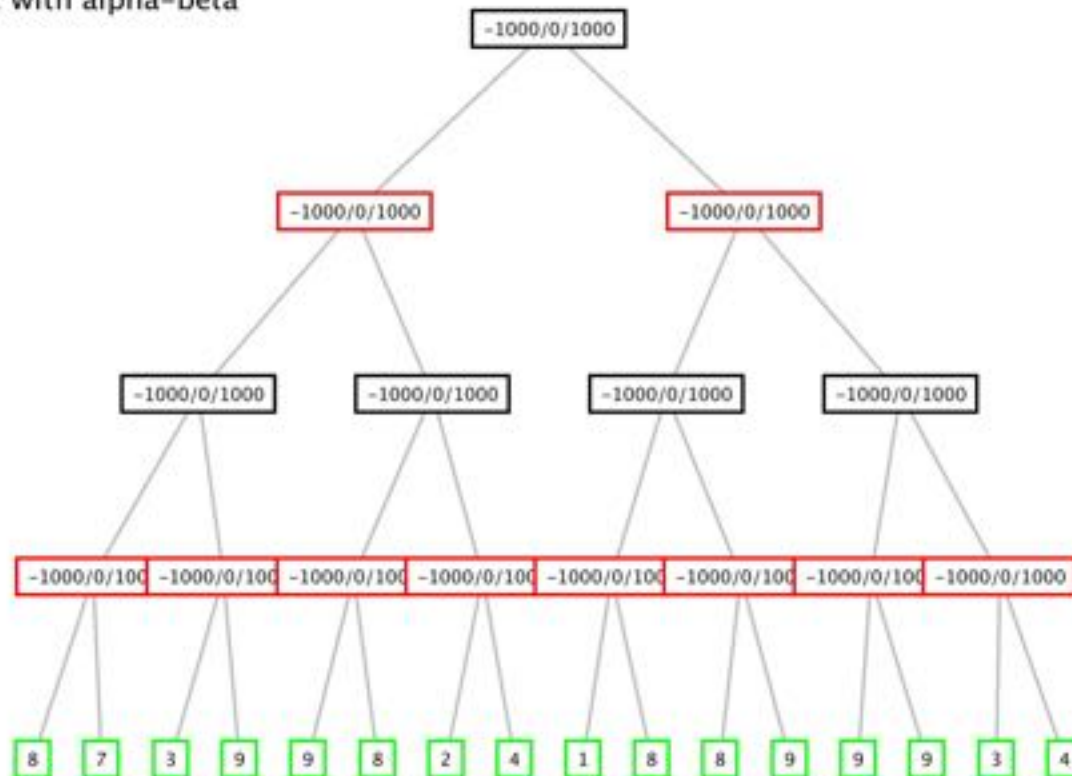
Theoretical minimum / actual / maximum: 7 / 0 / 16
Minimax with alpha-beta

MAX

MIN

MAX

MIN



Example

- ☒ Classroom example
- ☐ Quiz example
- ☐ Good news example

Arrangement

- ☐ Random numbers
- ☐ Increasing left to right
- ☐ Decreasing left to right

Search type

- ☐ Mini-max only
- ☒ Mini-max with alpha-beta

Shape

- ☒ 2 x 4
- ☐ 4 x 2
- ☐ 2 x 6
- ☐ 6 x 2


Delay

0.0 0.2 0.4 0.6 0.8 1.0

Ordering of static values can greatly affect alpha-beta pruning

- If the most favorable successor nodes for both MAX and MIN are on the LEFT so we explore them FIRST, then this leads to maximal pruning
- If the most favorable successor nodes for both MAX and MIN are on the RIGHT so they are explored LAST, then there is less pruning, possibly none at all
- Maximal pruning (in terms of branching factor b and tree depth d is approx: $\frac{2}{b^{d/2}}$, so can search down 2x as far using α - β search in this “good news” optimal case, compared to std minimax

Prune whenever $\alpha \geq \beta$

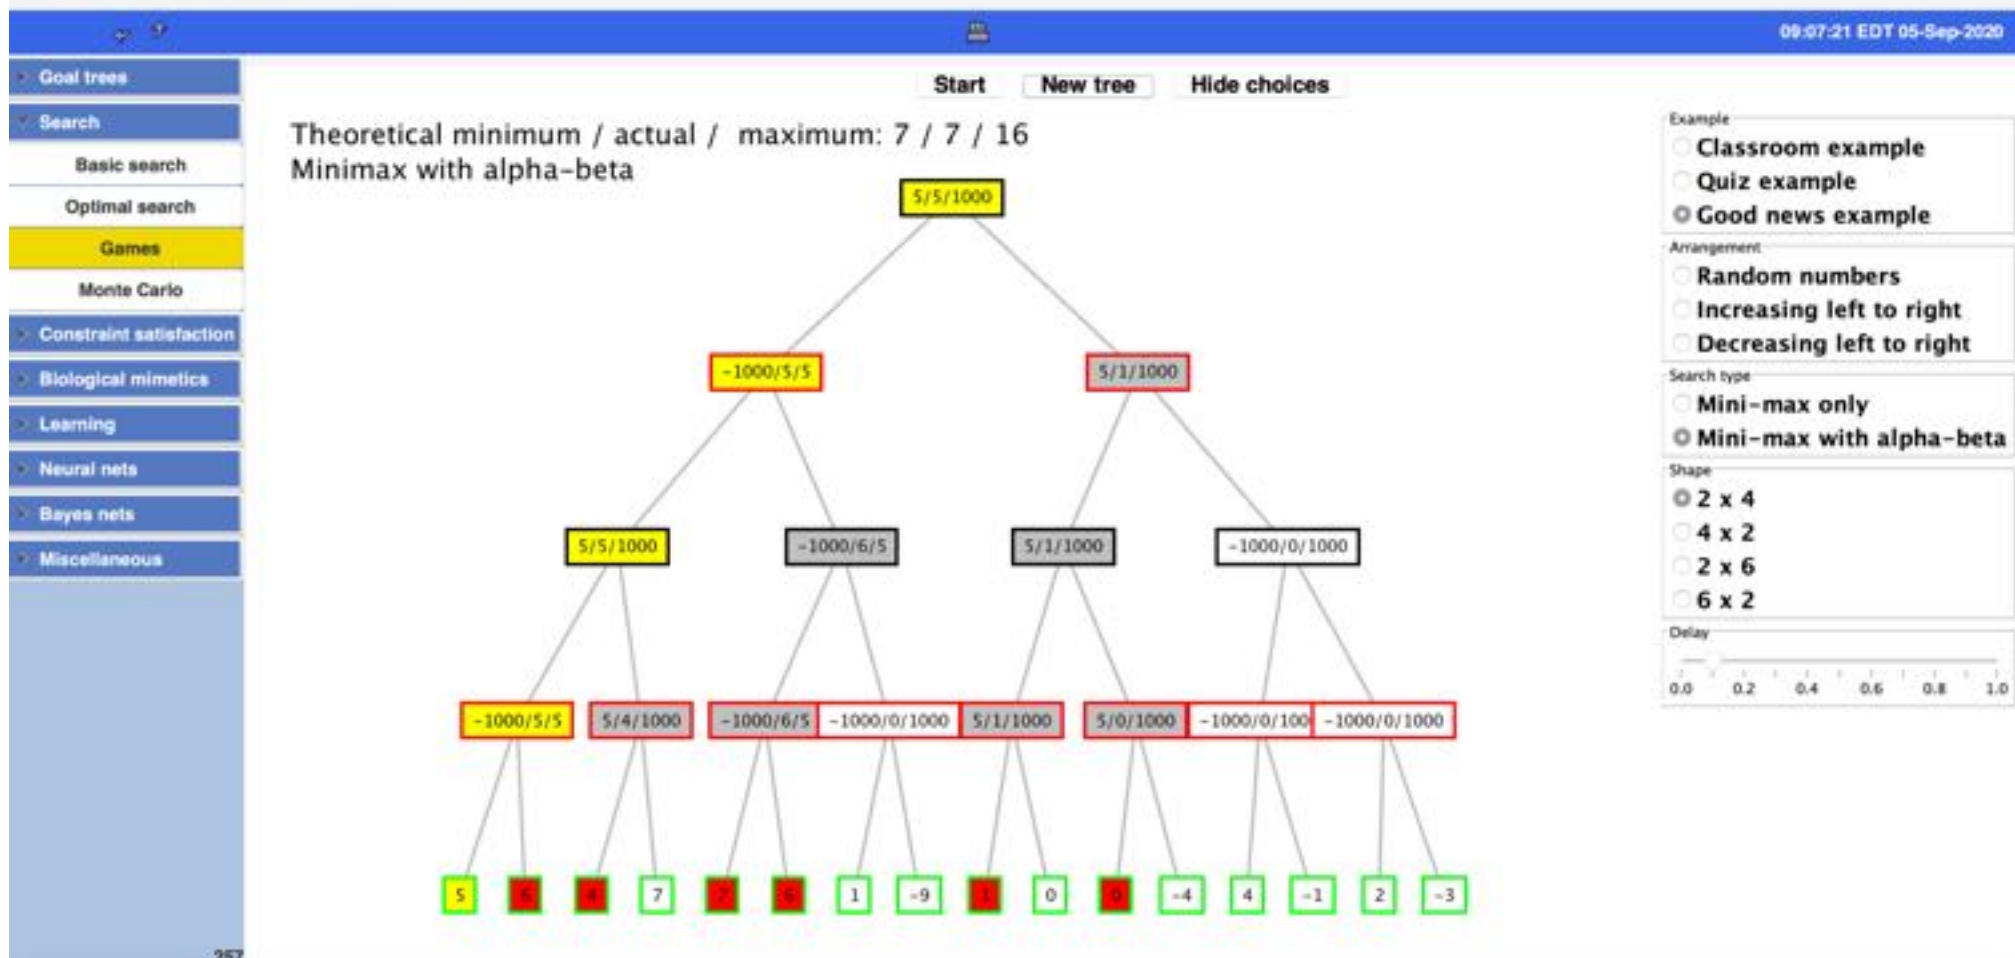


What would happen to pruning if nodes 3, 9 reversed? 3 Not pruned!

Optimal (good news) α - β game tree pruning



Games



$\approx 2b^{d/2}$ \approx maximum savings using α - β if
optimally ordered game tree

But suppose we run out of compute time?

The branching factor depends on the game &
board state – won't know for sure how deep
we can go in 1 minute....

Suppose we are playing blitz chess?

Will we always have a (good) move at hand?

Let's combine this thought with the idea of
optimally ordering the node evaluations...




Anytime algorithm




Martial arts principle

How not to run out of time



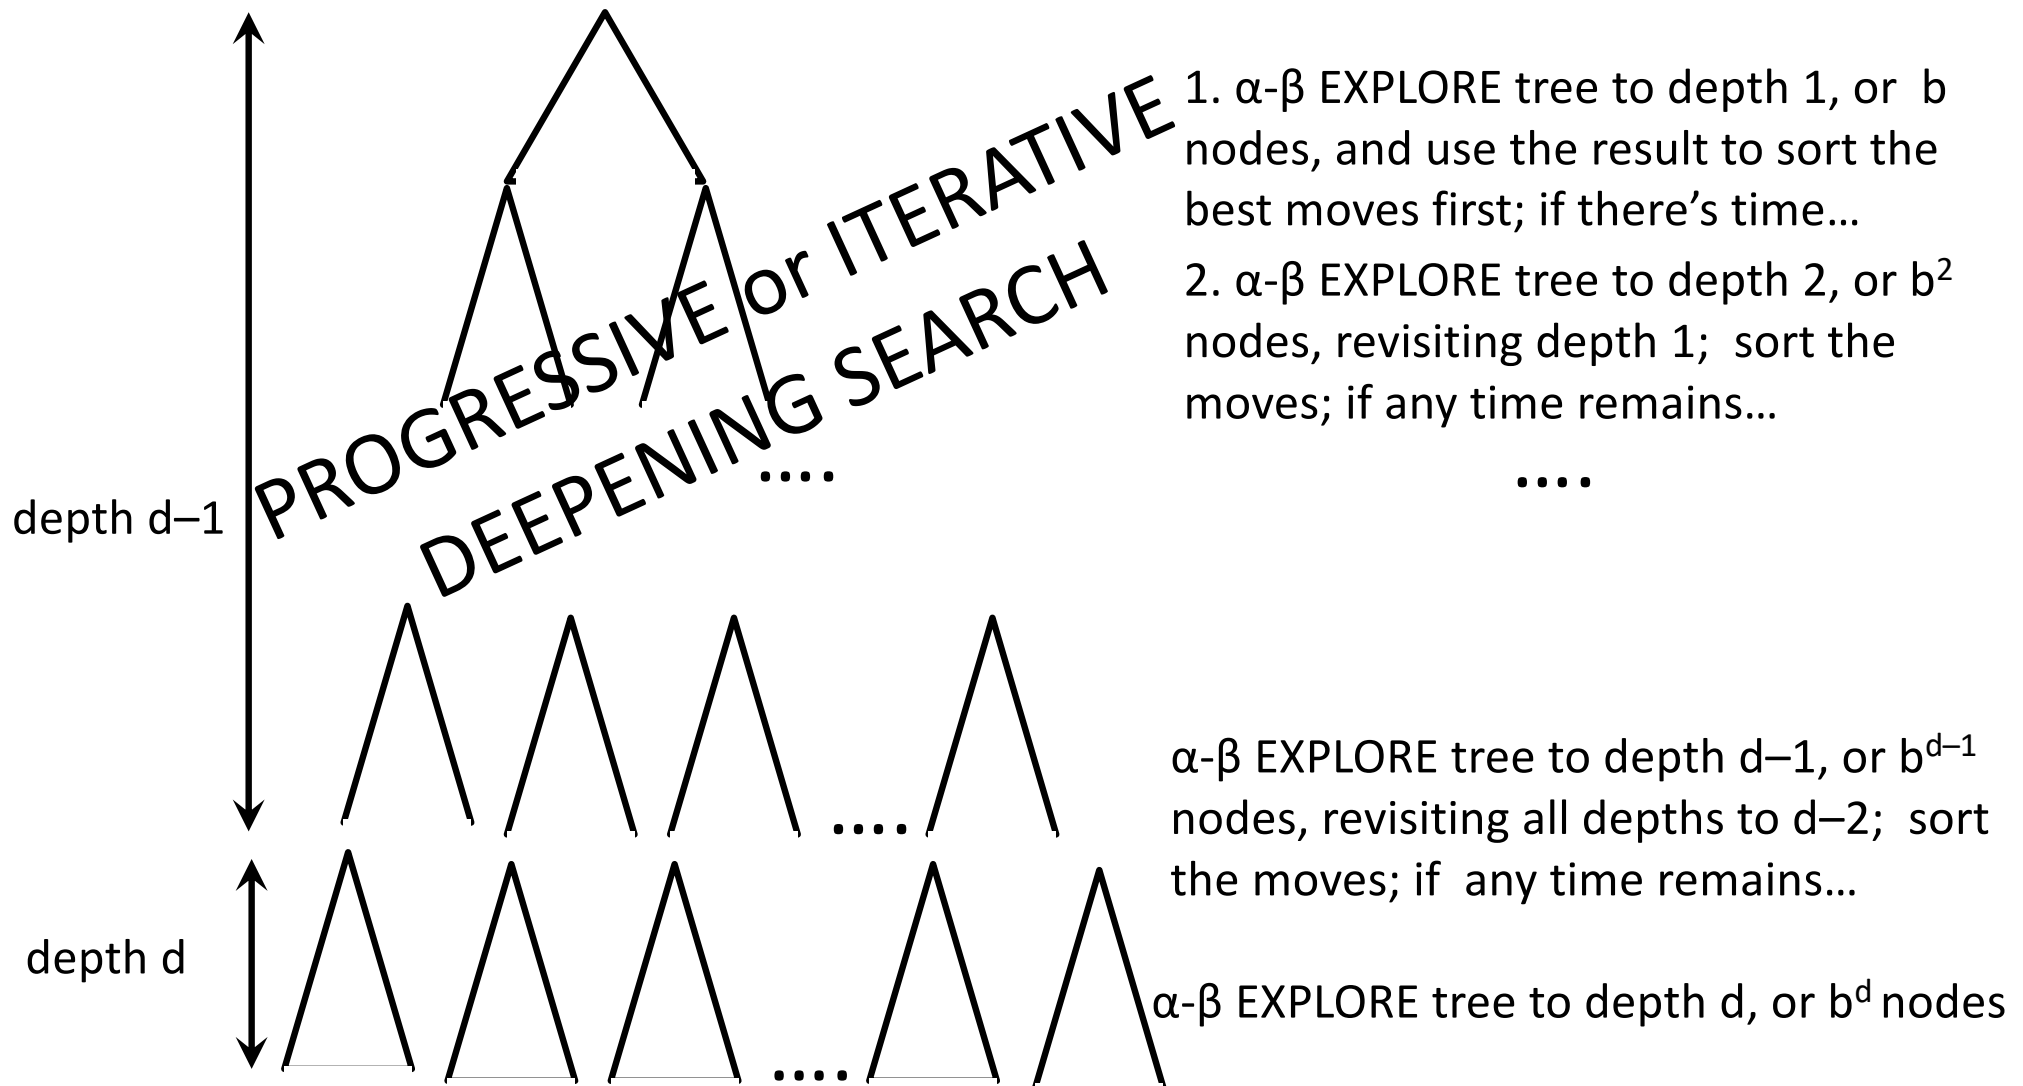
....



So, we can have a move in hand at level d-1 for only $1/b^{\text{th}}$ of computation cost required to go all the way down to level d

We can do this recursively to great benefit...

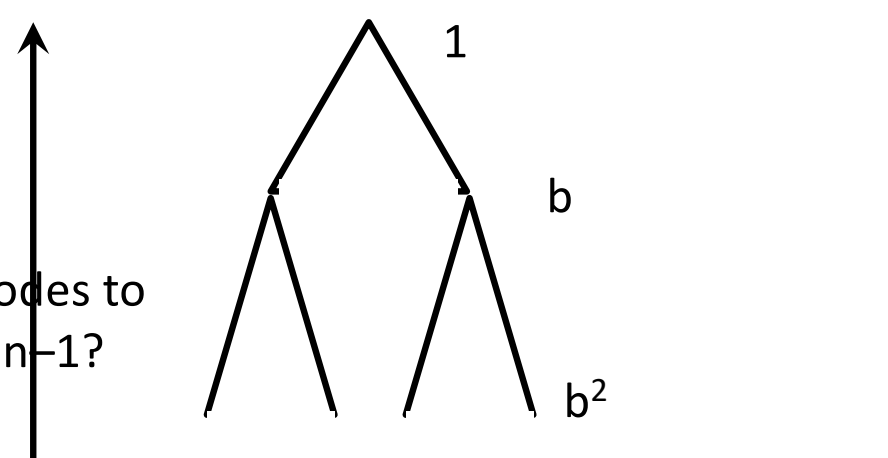
How not to run out of time (ignoring α - β for now)



Computational cost of progressive deepening?

How many total nodes S in game tree to depth $d-1$?


all nodes to depth $n-1$?



.....

$$S = 1 + b + b^2 + \dots + b^{d-1}$$
$$bS = b + b^2 + \dots + b^d$$
$$\Rightarrow bS - S = b^d - 1$$
$$S(b-1) = b^d - 1$$
$$S = \frac{b^d - 1}{b - 1} \approx b^{d-1}$$

depth d
nodes = b^d



..... b^{d-1}

So, search through all the game tree up thru depth $d-1$ is only time b^{d-1} , or $1/b^{\text{th}}$ of time to search to full depth d

Progressive (Iterative) deepening search

- Optimal DFS method (redundant use of space, but doesn't lose on arbitrary depth d trees)
- Earlier searches tend to improve the commonly used heuristics, so that a more accurate estimate of the score of various nodes at the final depth search can occur
- Because early iterations use small values for d they execute extremely quickly. This allows the algorithm to supply early indications of the result almost immediately, followed by refinements as d increases

Deep Blue,
IBM, circa
1997.
 $\approx 10^8$ static
evaluations
per sec



Deep Blue =
minimax + alpha-
beta +
progressive
deepening+
parallel
computing +
opening book +
end game +
secret sauce =
uneven tree
development

Gary Kasparov, then current world champion,
vs. Deep Blue, 1997



Is this intelligence?
Is this human intelligence?

Donald Byrne vs Bobby Fischer (black)

1956, move 17, Queen sacrifice



Gold star ideas today

- ★ Dead horse principle, aka “AWP”: α - β search
- ★ Martial arts principle – use adversary’s strength against them: progressive deepening
- ★ Anytime algorithms: progressive deepening
- ★ Simple \neq Trivial: sometimes, bulldozers work